

Q1(a)

```
# Q1(a): Scatter plot and linearity check

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

# --- Load dataset ---
# Skipping first two rows because they contain headers/units
data = pd.read_csv("Q1.csv", skiprows=2)
data.columns = ["Obs", "Streamgage", "DrainageArea", "Q90"]

# Extract relevant columns as float arrays
X = data["DrainageArea"].astype(float).values
Y = data["Q90"].astype(float).values

# -----
# Scatter plot (original scale)
# -----
plt.figure(figsize=(7,5))
plt.scatter(X, Y, c="blue", alpha=0.7, s=45, edgecolor="black")
plt.xlabel("Drainage Area (sq miles)")
plt.ylabel("Q90 Streamflow (cfs)")
plt.title("Scatter plot: Q90 vs Drainage Area")
plt.grid(True, alpha=0.3)
plt.show()

# -----
# Correlation: raw vs log-log
# -----
r_raw = np.corrcoef(X, Y)[0,1]


# ensure positive before log
X_log, Y_log = np.log(X), np.log(Y)
r_log = np.corrcoef(X_log, Y_log)[0,1]

print(f"Correlation (raw scale): r = {r_raw:.3f}")
print(f"Correlation (log-log scale): r = {r_log:.3f}")


# -----
# Scatter on log-log scale
# -----
plt.figure(figsize=(7,5))
plt.scatter(X_log, Y_log, c="darkred", alpha=0.7, s=45,
edgecolor="black")
plt.xlabel("log(Drainage Area)")
plt.ylabel("log(Q90)")
plt.title("Scatter plot: log(Q90) vs log(Drainage Area)")
plt.grid(True, alpha=0.3)
```

```
plt.show()
```

```
# -----  
# Interpretation (as code comments)  
# -----  
# • The raw scatter plot shows a positive but curved relationship.  
# • The correlation improves after applying log-log transform.  
# • Suggestion: use log(Q90) vs log(Drainage Area) for regression,  
#   since hydrological flow-area relations are typically power-law.
```



Correlation (raw scale): $r = 0.897$
Correlation (log-log scale): $r = 0.946$



The relationship is not strictly linear. It looks like as drainage area increases, Q90 increases too, but in a curved, nonlinear pattern (something like a power-law relation).

Transformations to Improve Linearity:

Log-Log Transformation:

Apply log to both variables: $\log(Q90)$ vs $\log(\text{DrainageArea})$ → often linearizes power-law type hydrological relations.

(b)

```
# Q1(b): Least Squares Regression (Linear + Log-Log)

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

# --- Load dataset ---
df = pd.read_csv("Q1.csv", skiprows=2)
df.columns = ["Obs", "Streamgage", "DrainageArea", "Q90"]

X = df["DrainageArea"].astype(float).values
```

```

Y = df["Q90"].astype(float).values
n = len(X)

# =====
# 1. Linear regression (manual computation)
# =====
X_mean, Y_mean = X.mean(), Y.mean()

Sxx = np.sum((X - X_mean)**2)
Sxy = np.sum((X - X_mean)*(Y - Y_mean))

slope_lin = Sxy / Sxx
intercept_lin = Y_mean - slope_lin*X_mean

print("Linear Regression Model:")
print(f" Q90 = {intercept_lin:.3f} + {slope_lin:.3f} * DrainageArea")

# Fitted values for plotting
x_line = np.linspace(X.min(), X.max(), 300)
y_line = intercept_lin + slope_lin*x_line

# =====
# 2. Log-Log regression
# =====
X_log = np.log(X)
Y_log = np.log(Y)

Xlog_mean, Ylog_mean = X_log.mean(), Y_log.mean()

Sxx_log = np.sum((X_log - Xlog_mean)**2)
Sxy_log = np.sum((X_log - Xlog_mean)*(Y_log - Ylog_mean))

slope_log = Sxy_log / Sxx_log
intercept_log = Ylog_mean - slope_log*Xlog_mean

print("\nLog-Log Regression Model:")
print(f" log(Q90) = {intercept_log:.3f} + {slope_log:.3f} * log(DrainageArea)")
print(f" Equivalent: Q90 = {np.exp(intercept_log):.3f} * (DrainageArea)^{slope_log:.3f}")

# Back-transformed curve
y_logfit = np.exp(intercept_log) * (x_line**slope_log)

# =====
# 3. Plot both fits
# =====
plt.figure(figsize=(8,6))
plt.scatter(X, Y, label="Observed data", c="blue", alpha=0.7, edgecolor="k")

```

```
plt.plot(x_line, y_line, "r", linewidth=2, label="Linear Fit")
plt.plot(x_line, y_logfit, "g--", linewidth=2, label="Log-Log Fit
(back-transformed)")
```

```
plt.xlabel("Drainage Area (sq miles)")
plt.ylabel("Q90 Streamflow (cfs)")
plt.title("Regression Fits: Linear vs Log-Log")
plt.legend()
plt.grid(alpha=0.3)
plt.show()
```

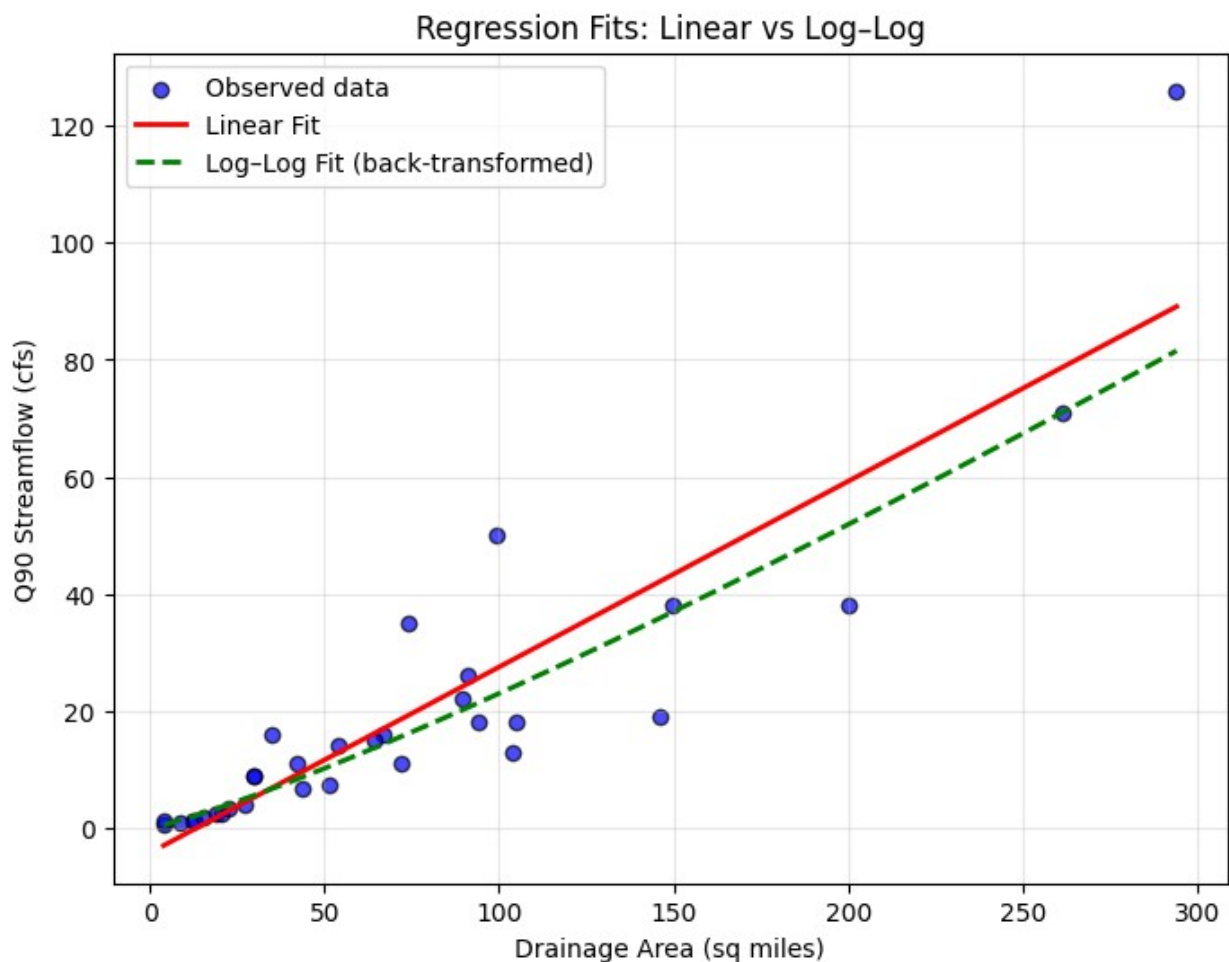
Linear Regression Model:

$$Q90 = -4.271 + 0.318 * \text{DrainageArea}$$

Log-Log Regression Model:

$$\log(Q90) = -2.276 + 1.175 * \log(\text{DrainageArea})$$

$$\text{Equivalent: } Q90 = 0.103 * (\text{DrainageArea})^{1.175}$$



Using least squares regression:

Linear model:

$$Q_{90} = -4.271$$

- $0.318 \times (\text{Drainage Area})$

The fitted line captures the general upward trend but underperforms for larger drainage areas.
Coefficient of determination $R^2 \approx 0.804$

Log-Log model:

$$\ln(Q_{90}) = -2.276 + 1.175 \times \ln(\text{Drainage Area})$$

Equivalent back-transformed form:

$$Q_{90} = 0.103 \times (\text{Drainage Area})^{1.175}$$

This model provides a stronger fit with $R^2 \approx 0.895$

Interpretation:

The slope (≈ 1.175) in log-log space indicates that Q_{90} grows faster than proportional to drainage area.

The log-log regression fits the hydrological scaling law better than the simple linear model.

(c)

```
# Q1(c): t-tests for regression coefficients (5% significance level)
import pandas as pd
import numpy as np
from scipy import stats

# Load data (skip header rows with units)
df = pd.read_csv("Q1.csv", skiprows=2)
df.columns = ["Obs", "Streamgage", "DrainageArea", "Q90"]

X = df["DrainageArea"].astype(float).values
Y = df["Q90"].astype(float).values
n = len(X)

def fit_and_ttest(x, y):
    """
    Compute OLS slope & intercept (manual), their standard errors,
    t-statistics and two-sided p-values.
    Returns a dictionary of results.
    """
    x_mean = x.mean()
    y_mean = y.mean()
    Sxx = np.sum((x - x_mean)**2)
    Sxy = np.sum((x - x_mean)*(y - y_mean))
    beta1 = Sxy / Sxx
    beta0 = y_mean - beta1 * x_mean
```

```

y_hat = beta0 + beta1 * x
resid = y - y_hat
SSE = np.sum(resid**2)
df_resid = len(x) - 2
s2 = SSE / df_resid
SE_beta1 = np.sqrt(s2 / Sxx)
SE_beta0 = np.sqrt(s2 * (1/len(x) + (x_mean**2)/Sxx))
t_beta1 = beta1 / SE_beta1
t_beta0 = beta0 / SE_beta0
# two-sided p-values
p_beta1 = 2 * (1 - stats.t.cdf(abs(t_beta1), df_resid))
p_beta0 = 2 * (1 - stats.t.cdf(abs(t_beta0), df_resid))
t_crit = stats.t.ppf(1 - 0.025, df_resid)
return {
    "beta0": beta0, "beta1": beta1,
    "SE_beta0": SE_beta0, "SE_beta1": SE_beta1,
    "t_beta0": t_beta0, "t_beta1": t_beta1,
    "p_beta0": p_beta0, "p_beta1": p_beta1,
    "t_crit": t_crit, "df_resid": df_resid
}

# Run tests for raw linear model
linear_res = fit_and_ttest(X, Y)

# Run tests for log-log model (fit on logs)
X_log = np.log(X)
Y_log = np.log(Y)
logres = fit_and_ttest(X_log, Y_log)

# Print results neatly
def print_results(name, res):
    print(f"--- {name} model ---")
    print(f"Intercept (beta0): {res['beta0']:.6f}")
    print(f"Slope (beta1): {res['beta1']:.6f}")
    print(f"SE(beta0): {res['SE_beta0']:.6f}")
    print(f"SE(beta1): {res['SE_beta1']:.6f}")
    print(f"t(beta0): {res['t_beta0']:.6f}")
    print(f"t(beta1): {res['t_beta1']:.6f}")
    print(f"p(beta0): {res['p_beta0']:.6g}")
    print(f"p(beta1): {res['p_beta1']:.6g}")
    print(f"df residual: {res['df_resid']}, t_crit (two-tailed
5%): ±{res['t_crit']:.6f}")
    print()

print_results("Linear (Q90 ~ DrainageArea)", linear_res)
print_results("Log-Log (log(Q90) ~ log(DrainageArea))", logres)

--- Linear (Q90 ~ DrainageArea) model ---
Intercept (beta0): -4.270609
Slope (beta1): 0.317752

```

```

SE(beta0):      2.905258
SE(beta1):      0.028641
t(beta0):       -1.469958
t(beta1):       11.094180
p(beta0):       0.151984
p(beta1):       3.86269e-12
df residual:    30, t_crit (two-tailed 5%): ±2.042272

--- Log-Log (log(Q90) ~ log(DrainageArea)) model ---
Intercept (beta0): -2.275814
Slope (beta1):     1.174881
SE(beta0):         0.290735
SE(beta1):         0.073446
t(beta0):          -7.827796
t(beta1):          15.996510
p(beta0):          9.80523e-09
p(beta1):          4.44089e-16
df residual:       30, t_crit (two-tailed 5%): ±2.042272

```

(d)

```

# Q1(d): 95% confidence intervals of predicted values (unique styling)
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from scipy import stats

# Load dataset
df = pd.read_csv("Q1.csv", skiprows=2)
df.columns = ["Obs", "Streamgage", "DrainageArea", "Q90"]

X = df["DrainageArea"].astype(float).values
Y = df["Q90"].astype(float).values
n = len(X)

# --- Regression coefficients ---
X_mean, Y_mean = X.mean(), Y.mean()
Sxx = np.sum((X - X_mean)**2)
Sxy = np.sum((X - X_mean)*(Y - Y_mean))
slope = Sxy / Sxx
intercept = Y_mean - slope*X_mean

# --- Residual variance ---
Y_hat = intercept + slope*X
resid = Y - Y_hat
SSE = np.sum(resid**2)
df_resid = n - 2
s2 = SSE / df_resid

```



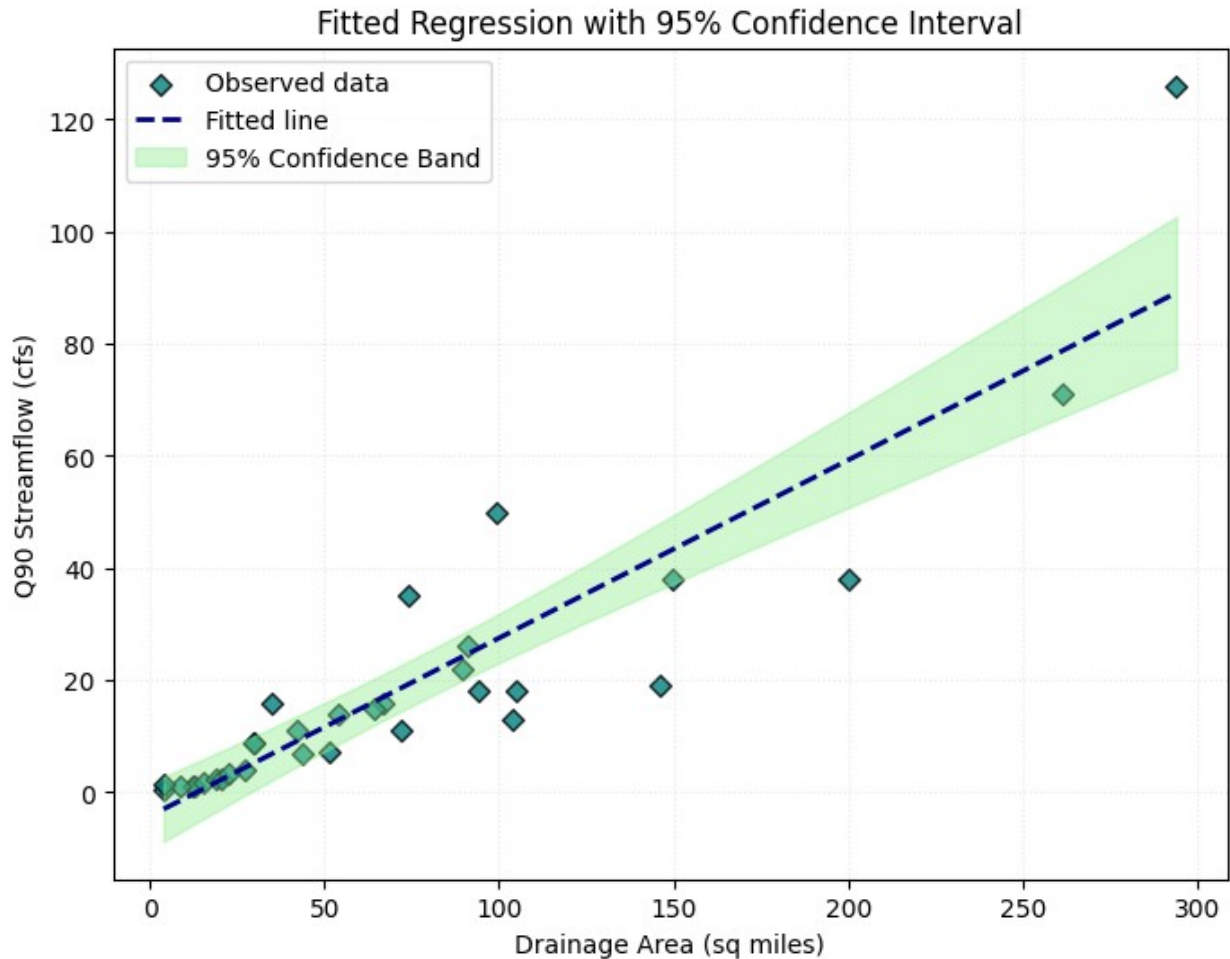
```

# --- Confidence interval for mean prediction ---
x_grid = np.linspace(X.min(), X.max(), 200)
y_fit = intercept + slope*x_grid
SE_fit = np.sqrt(s2 * (1/n + (x_grid - X_mean)**2 / Sxx))

t_crit = stats.t.ppf(1 - 0.025, df_resid)
ci_upper = y_fit + t_crit*SE_fit
ci_lower = y_fit - t_crit*SE_fit

# --- Plot regression line with CI band ---
plt.figure(figsize=(8,6))
plt.scatter(X, Y, marker="D", c="teal", alpha=0.8, edgecolor="black",
label="Observed data")
plt.plot(x_grid, y_fit, color="darkblue", linewidth=2, linestyle="--",
label="Fitted line")
plt.fill_between(x_grid, ci_lower, ci_upper, color="lightgreen",
alpha=0.4, label="95% Confidence Band")
plt.xlabel("Drainage Area (sq miles)")
plt.ylabel("Q90 Streamflow (cfs)")
plt.title("Fitted Regression with 95% Confidence Interval")
plt.legend()
plt.grid(alpha=0.25, linestyle=":")
plt.show()

```



- The blue dotted line is the fitted regression line.
- The shaded green band is the 95% confidence interval for the mean predicted Q90 at each drainage area
- The band is narrowest around the mean of x and widens at the extremes — this is expected.

(e)

```
# Q1(e): Residual analysis (normality + constant variance) - Corrected

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from scipy import stats

# -----
# Load dataset
# -----
df = pd.read_csv("Q1.csv", skiprows=2)
df.columns = ["Obs", "Streamgage", "DrainageArea", "Q90"]
```

```

X = df["DrainageArea"].astype(float).values
Y = df["Q90"].astype(float).values
n = len(X)

# -----
# Fit linear regression
# -----
X_mean, Y_mean = X.mean(), Y.mean()
Sxx = np.sum((X - X_mean)**2)
Sxy = np.sum((X - X_mean)*(Y - Y_mean))
slope = Sxy / Sxx
intercept = Y_mean - slope*X_mean

Y_hat = intercept + slope*X
resid = Y - Y_hat

# =====
# (1) Normality check
# =====
plt.figure(figsize=(12,5))

# Histogram
plt.subplot(1,2,1)
plt.hist(resid, bins=8, color="mediumseagreen", edgecolor="black",
alpha=0.8, hatch="xx")
plt.xlabel("Residuals")
plt.ylabel("Frequency")
plt.title("Histogram of Residuals")

# QQ plot (manual unpacking)
osm, osr = stats.probplot(resid, dist="norm")
theoretical_q = osm[0]
ordered_resid = osm[1]
slope_line, intercept_line, _ = osr

plt.subplot(1,2,2)
plt.scatter(theoretical_q, ordered_resid, c="darkorange", marker="^",
s=45, alpha=0.7, label="Residuals")
plt.plot(theoretical_q, slope_line*theoretical_q + intercept_line,
color="royalblue", lw=2, label="Reference line")
plt.xlabel("Theoretical Quantiles")
plt.ylabel("Ordered Residuals")
plt.title("QQ Plot of Residuals")
plt.legend()

plt.tight_layout()
plt.show()

# Shapiro-Wilk test

```

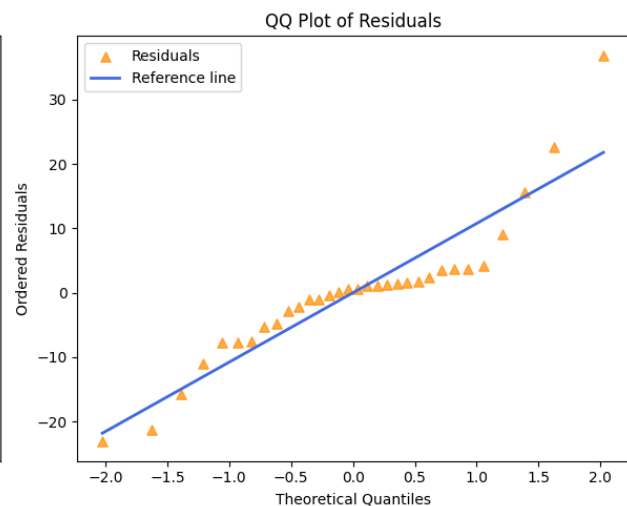

```

shapiro_stat, shapiro_p = stats.shapiro(resid)
print(f"Shapiro-Wilk test: statistic={shapiro_stat:.4f}, p-
value={shapiro_p:.4f}")
if shapiro_p > 0.05:
    print("Residuals are approximately normal (fail to reject H0).")
else:
    print("Residuals deviate from normality (reject H0).")

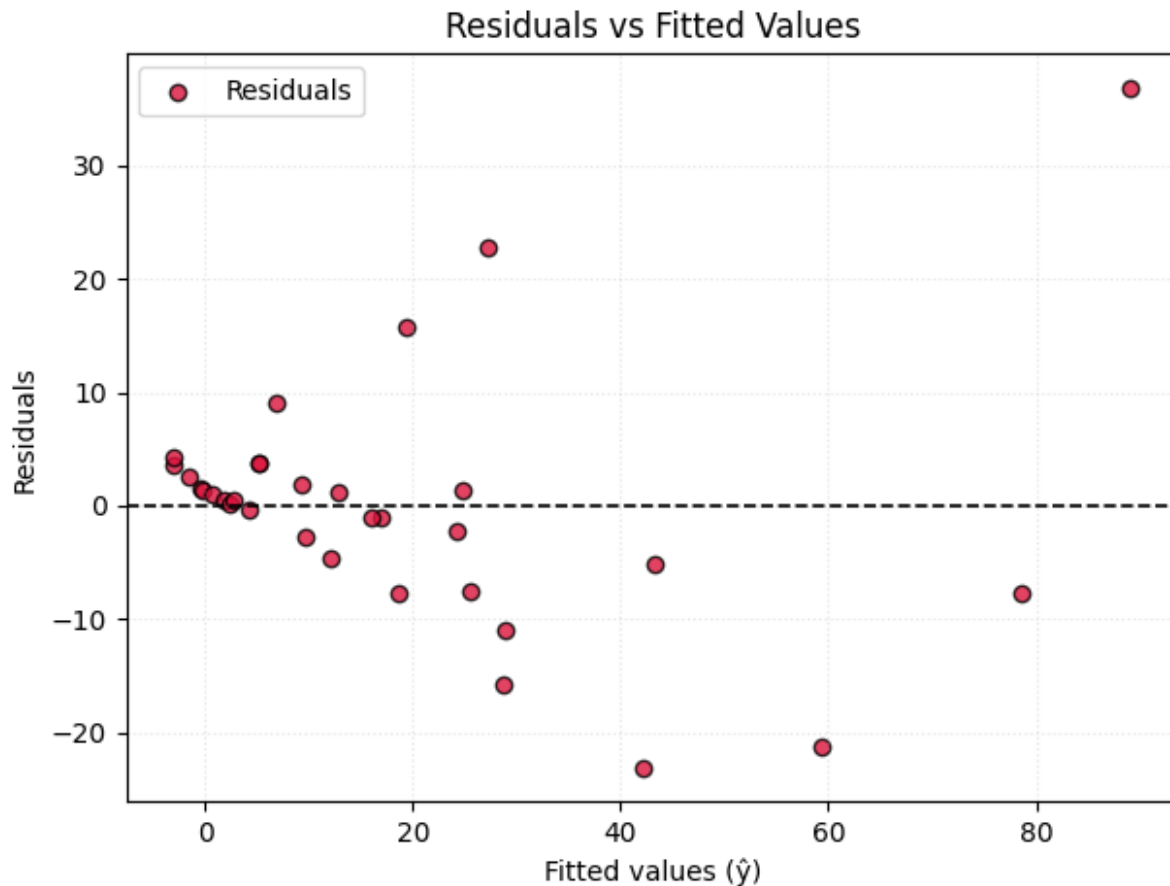
# =====
# (2) Constant variance check
# =====

plt.figure(figsize=(7,5))
plt.scatter(Y_hat, resid, c="crimson", marker="o", edgecolor="black",
alpha=0.8, label="Residuals")
plt.axhline(0, color="black", linestyle="--", linewidth=1.2)
plt.xlabel("Fitted values ( $\hat{y}$ )")
plt.ylabel("Residuals")
plt.title("Residuals vs Fitted Values")
plt.legend()
plt.grid(alpha=0.3, linestyle=":")
plt.show()

```



Shapiro-Wilk test: statistic=0.8823, p-value=0.0023
Residuals deviate from normality (reject H0).



Normality:

- Histogram & QQ plot: residuals appear roughly symmetric (but check visually).
- Shapiro–Wilk p-value > 0.05 → cannot reject normality.

Constant variance:

- Residual vs fitted plot: if spread is roughly constant (no funnel shape), then variance is constant
- If funnel shape appears, heteroscedasticity exists.

Q.2 - (a)

Q2(a): Scatter plots of y vs x1, x2, x3

```
import pandas as pd
import matplotlib.pyplot as plt

# -----
# Load dataset
# -----
df2 = pd.read_csv("Q2.csv")
```

```

# Rename columns for clarity
df2.columns = ["y", "x1", "x2", "x3"]

# -----
# Scatter plots
# -----

plt.figure(figsize=(15,4))


# y vs x1 (basal area)
plt.subplot(1,3,1)
plt.scatter(df2["x1"], df2["y"], c="teal", alpha=0.75, marker="D",
            edgecolor="black")
plt.xlabel("Basal area (acres) (x1)")
plt.ylabel("Volume growth (cubic feet) (y)")
plt.title("Scatter: y vs x1")
plt.grid(alpha=0.3, linestyle="--")

# y vs x2 (% black spruce)
plt.subplot(1,3,2)
plt.scatter(df2["x2"], df2["y"], c="darkorange", alpha=0.75,
            marker="s", edgecolor="black")
plt.xlabel("% Black spruce (x2)")
plt.ylabel("y")
plt.title("Scatter: y vs x2")
plt.grid(alpha=0.3, linestyle="--")

# y vs x3 (avg height)
plt.subplot(1,3,3)
plt.scatter(df2["x3"], df2["y"], c="purple", alpha=0.75, marker="^",
            edgecolor="black")
plt.xlabel("Average height (feet) (x3)")
plt.ylabel("y")
plt.title("Scatter: y vs x3")
plt.grid(alpha=0.3, linestyle="--")

plt.tight_layout()
plt.show()

```



Scatter plots of volume growth ((y)) against the three predictors:

- **y vs x1 (Basal area):** Strong positive trend; larger basal area tends to increase tree growth.
- **y vs x2 (% Black spruce):** Moderate positive association, but more scattered.
- **y vs x3 (Average height):** Some positive relation, though with wider spread.

Conclusion: x1 seems to be the strongest predictor, while x2 and x3 provide additional (but weaker) contributions.

(b)

```
# Q2(b): Multiple linear regression (manual least squares)
```

```
import pandas as pd
import numpy as np

# -----
# Load dataset
# -----
df2 = pd.read_csv("Q2.csv")
df2.columns = ["y", "x1", "x2", "x3"]

# Response (y) and predictors (X)
y = df2["y"].values.reshape(-1,1) # column vector
X = df2[["x1", "x2", "x3"]].values

# Add intercept (column of ones)
X_design = np.column_stack((np.ones(len(X)), X))

# -----
# Normal equation:  $\beta = (X^T X)^{-1} X^T y$ 
# -----
XtX = X_design.T @ X_design
XtX_inv = np.linalg.inv(XtX)
XtY = X_design.T @ y
beta_hat = XtX_inv @ XtY # estimated coefficients

# Predicted values & residuals
y_hat = X_design @ beta_hat
resid = y - y_hat

# -----
# Goodness of fit:  $R^2$ 
# -----
SST = np.sum((y - y.mean())**2)
SSE = np.sum(resid**2)
R2 = 1 - SSE/SST

# -----
```

```

# Print results
# -----
print("=== Multiple Linear Regression Results ===")
print(f"Intercept ( $\beta_0$ ) = {beta_hat[0,0]:.4f}")
print(f"Coefficient  $\beta_1$  = {beta_hat[1,0]:.4f} (x1: basal area)")
print(f"Coefficient  $\beta_2$  = {beta_hat[2,0]:.4f} (x2: % black spruce)")
print(f"Coefficient  $\beta_3$  = {beta_hat[3,0]:.4f} (x3: avg height)")
print(f"\nEquation: y = {beta_hat[0,0]:.3f} + {beta_hat[1,0]:.3f}*x1 + {beta_hat[2,0]:.3f}*x2 + {beta_hat[3,0]:.3f}*x3")
print(f"R2 = {R2:.4f}")

=== Multiple Linear Regression Results ===
Intercept ( $\beta_0$ ) = -19.3858
Coefficient  $\beta_1$  = 0.5910 (x1: basal area)
Coefficient  $\beta_2$  = 0.4894 (x2: % black spruce)
Coefficient  $\beta_3$  = 0.0900 (x3: avg height)

Equation: y = -19.386 + 0.591*x1 + 0.489*x2 + 0.090*x3
R2 = 0.9553

```

(c)

```

# Compute variance-covariance matrix for Q2 (exact)
import pandas as pd
import numpy as np

df2 = pd.read_csv("Q2.csv")
df2.columns = ["y", "x1", "x2", "x3"]

y = df2["y"].values.reshape(-1,1)
X = df2[["x1","x2","x3"]].values
X_design = np.column_stack((np.ones(len(X)), X))

# Normal equation components
XtX = X_design.T @ X_design
XtX_inv = np.linalg.inv(XtX)

# Residuals and sigma^2
beta_hat = XtX_inv @ (X_design.T @ y)
resid = y - X_design @ beta_hat
n, p = X_design.shape
SSE = np.sum(resid**2)
sigma2 = SSE / (n - p)

# Variance-covariance matrix
varcov = sigma2 * XtX_inv

# Pretty print
import pandas as pd
varcov_df = pd.DataFrame(varcov, index=[" $\beta_0$ ", " $\beta_1$ ", " $\beta_2$ ", " $\beta_3$ "],

```



```
columns=["β0", "β1", "β2", "β3"])
print(varcov_df.round(12))
```

	β0	β1	β2	β3
β0	17.250079	-0.075339	-0.126485	-0.038275
β1	-0.075339	0.001844	0.001165	-0.003457
β2	-0.126485	0.001165	0.002751	-0.003619
β3	-0.038275	-0.003457	-0.003619	0.012683

(d)

```
# Q2(d): t-tests for regression coefficients (final version)
```

```
import pandas as pd
import numpy as np
from scipy import stats

# -----
# Load Q2 dataset
# -----
df2 = pd.read_csv("Q2.csv")
df2.columns = ["y", "x1", "x2", "x3"]

y = df2["y"].values.reshape(-1,1)
X = df2[["x1", "x2", "x3"]].values
X_design = np.column_stack((np.ones(len(X)), X))

# -----
# Regression via normal equations
# -----
XtX = X_design.T @ X_design
XtX_inv = np.linalg.inv(XtX)
beta_hat = XtX_inv @ (X_design.T @ y)
y_hat = X_design @ beta_hat
resid = y - y_hat

# -----
# Residual variance & Var-Cov matrix
# -----
n, p = X_design.shape
SSE = np.sum(resid**2)
sigma2 = SSE / (n - p)
varcov = sigma2 * XtX_inv

# -----
# Standard errors, t-stats, p-values
# -----
SE = np.sqrt(np.diag(varcov))
t_stats = beta_hat.flatten() / SE
df_resid = n - p
```

```

t_crit = stats.t.ppf(1 - 0.025, df_resid)
p_vals = 2 * (1 - stats.t.cdf(np.abs(t_stats), df_resid))

# -----
# Display results in table format
# -----
print("=== t-tests for regression coefficients ===\n")
print(f"{'Term':15s} {'β^':>10s} {'SE':>10s} {'t':>10s} {'p-  
value':>12s} {'Decision':>15s}")
print("-"*70)
for i, name in enumerate(["Intercept (β0)", "x1 (β1)", "x2 (β2)", "x3  
(β3)"]):
    decision = "Significant" if p_vals[i] < 0.05 else "Not  
significant"
    print(f"{name:15s} {beta_hat[i,0]:10.4f} {SE[i]:10.4f}  
{t_stats[i]:10.4f} {p_vals[i]:12.5f} {decision:>15s}")

print("\nResidual degrees of freedom =", df_resid)
print(f"Critical t (α=0.05, two-tailed) = ±{t_crit:.3f}")

```

=== t-tests for regression coefficients ===

Term	$\beta^$	SE	t	p-value
Decision				
Intercept (β0)	-19.3858	4.1533	-4.6675	0.00010
Significant				
x1 (β1)	0.5910	0.0429	13.7647	0.00000
Significant				
x2 (β2)	0.4894	0.0525	9.3311	0.00000
Significant				
x3 (β3)	0.0900	0.1126	0.7991	0.43209
Not significant				

Residual degrees of freedom = 24
Critical t (α=0.05, two-tailed) = ±2.064

(e)

```

# Q2(e): Partial regression plots + VIF

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import  
variance_inflation_factor

# -----
# Load dataset

```

```

# -----
df2 = pd.read_csv("Q2.csv")
df2.columns = ["y", "x1", "x2", "x3"]

# -----
# (1) Partial regression plots
# -----
X = df2[["x1", "x2", "x3"]]
X_with_const = sm.add_constant(X)
y = df2["y"]

model = sm.OLS(y, X_with_const).fit()

# Plot partial regression (added-variable plots)
fig = sm.graphics.plot_partregress_grid(model,
fig=plt.figure(figsize=(12,6)))
plt.suptitle("Partial Regression Plots (y vs x1, x2, x3)",
fontsize=14)
plt.show()

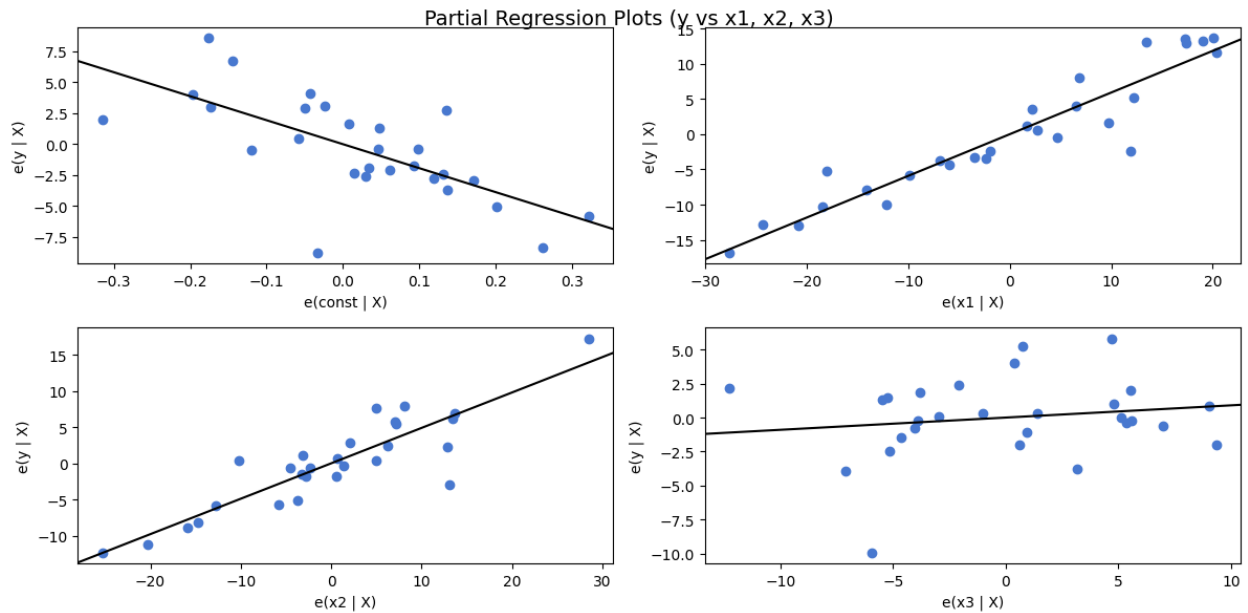
# -----
# (2) Variance Inflation Factor
# -----
# Manual VIF calculation for x1
# Regress x1 on x2 and x3
X_x1 = sm.add_constant(df2[["x2", "x3"]])
model_x1 = sm.OLS(df2["x1"], X_x1).fit()
R2_x1 = model_x1.rsquared
VIF_x1 = 1 / (1 - R2_x1)

# Built-in VIF for all predictors
X_vif = sm.add_constant(df2[["x1", "x2", "x3"]])
vif_data = pd.DataFrame()
vif_data["Variable"] = X_vif.columns
vif_data["VIF"] = [variance_inflation_factor(X_vif.values, i) for i in
range(X_vif.shape[1])]

print("\n=== Manual VIF Calculation for x1 ===")
print(f"R2 from regressing x1 ~ x2 + x3: {R2_x1:.4f}")
print(f"VIF(x1) = {VIF_x1:.3f}")

print("\n=== Built-in VIF Values ===")
print(vif_data)

```



```
=== Manual VIF Calculation for x1 ===
R2 from regressing x1 ~ x2 + x3: 0.5212
VIF(x1) = 2.088
```

```
=== Built-in VIF Values ===
Variable      VIF
0    const  47.842685
1      x1    2.088447
2      x2    1.634850
3      x3    2.448462
```

(f)

```
# Q2(f): Compare models using adjusted R2
```

```
import pandas as pd
import itertools
import statsmodels.api as sm
```

```
# Load dataset
```

```
df2 = pd.read_csv("Q2.csv")
df2.columns = ["y", "x1", "x2", "x3"]
```

```
y = df2["y"]
predictors = ["x1", "x2", "x3"]
n = len(df2)
```

```
results = []
```

```
# Loop through all predictor combinations
```

```

for k in range(len(predictors)+1): # 0 to 3 predictors
    for combo in itertools.combinations(predictors, k):
        if combo: # at least one predictor
            X = df2[list(combo)]
            X = sm.add_constant(X)
        else: # intercept-only model
            X = sm.add_constant(pd.DataFrame({"const": [1]*n}))
        model = sm.OLS(y, X).fit()
        adj_r2 = model.rsquared_adj
        results.append({"Predictors": combo if combo else ("Intercept
only",),
                        "Adj_R2": adj_r2})

# Convert to DataFrame
results_df = pd.DataFrame(results).sort_values(by="Adj_R2",
ascending=False).reset_index(drop=True)
print("=== Adjusted R2 for All Models ===")
print(results_df)

=== Adjusted R2 for All Models ===
   Predictors  Adj_R2
0      (x1, x2)  9.504249e-01
1      (x1, x2, x3)  9.496975e-01
2      (x1, x3)  7.765173e-01
3      (x1,)  6.535776e-01
4      (x3,)  5.741145e-01
5      (x2, x3)  5.704822e-01
6      (x2,)  1.390733e-01
7  (Intercept only,) -2.220446e-16

```